

# Convolutional Neural Networks

**Dr. Unchalisa Taetrugool**

Department of Computer Engineering, Faculty of Engineering  
King Mongkut's University of Technology Thonburi

# Convolutional Neural Network











Image source: "Unsupervised Learning of Hierarchical Representations with Convolutional Deep Belief Networks" ICML 2009 & Comm. ACM 2011.  
Honglak Lee, Roger Grosse, Rajesh Ranganath, and Andrew Ng.

2  NVIDIA

# Image Recognition: Simple Example




# Image Recognition: Simple Example



|    |    |
|----|----|
| 1  | -1 |
| -1 | 1  |

|   |    |    |   |
|---|----|----|---|
| 1 | -1 | -1 | 1 |
|---|----|----|---|


*convolve filter/filter map*



|    |    |
|----|----|
| -1 | 1  |
| 1  | -1 |


|    |   |   |    |
|----|---|---|----|
| -1 | 1 | 1 | -1 |
|----|---|---|----|

# Image Recognition: Simple Example



| + | +  | +  | + |
|---|----|----|---|
| 1 | -1 | -1 | 1 |


$$+1 \quad -1 \quad -1 \quad +1 = 0$$



| +  | + | + | +  |
|----|---|---|----|
| -1 | 1 | 1 | -1 |


$$-1 \quad 1 \quad 1 \quad -1 = 0$$

# Image Recognition: Simple Example



|    |    |
|----|----|
| 1  | -1 |
| -1 | 1  |

|   |    |    |   |
|---|----|----|---|
| 1 | -1 | -1 | 1 |
|---|----|----|---|



|    |    |
|----|----|
| -1 | 1  |
| 1  | -1 |

|    |   |   |    |
|----|---|---|----|
| -1 | 1 | 1 | -1 |
|----|---|---|----|

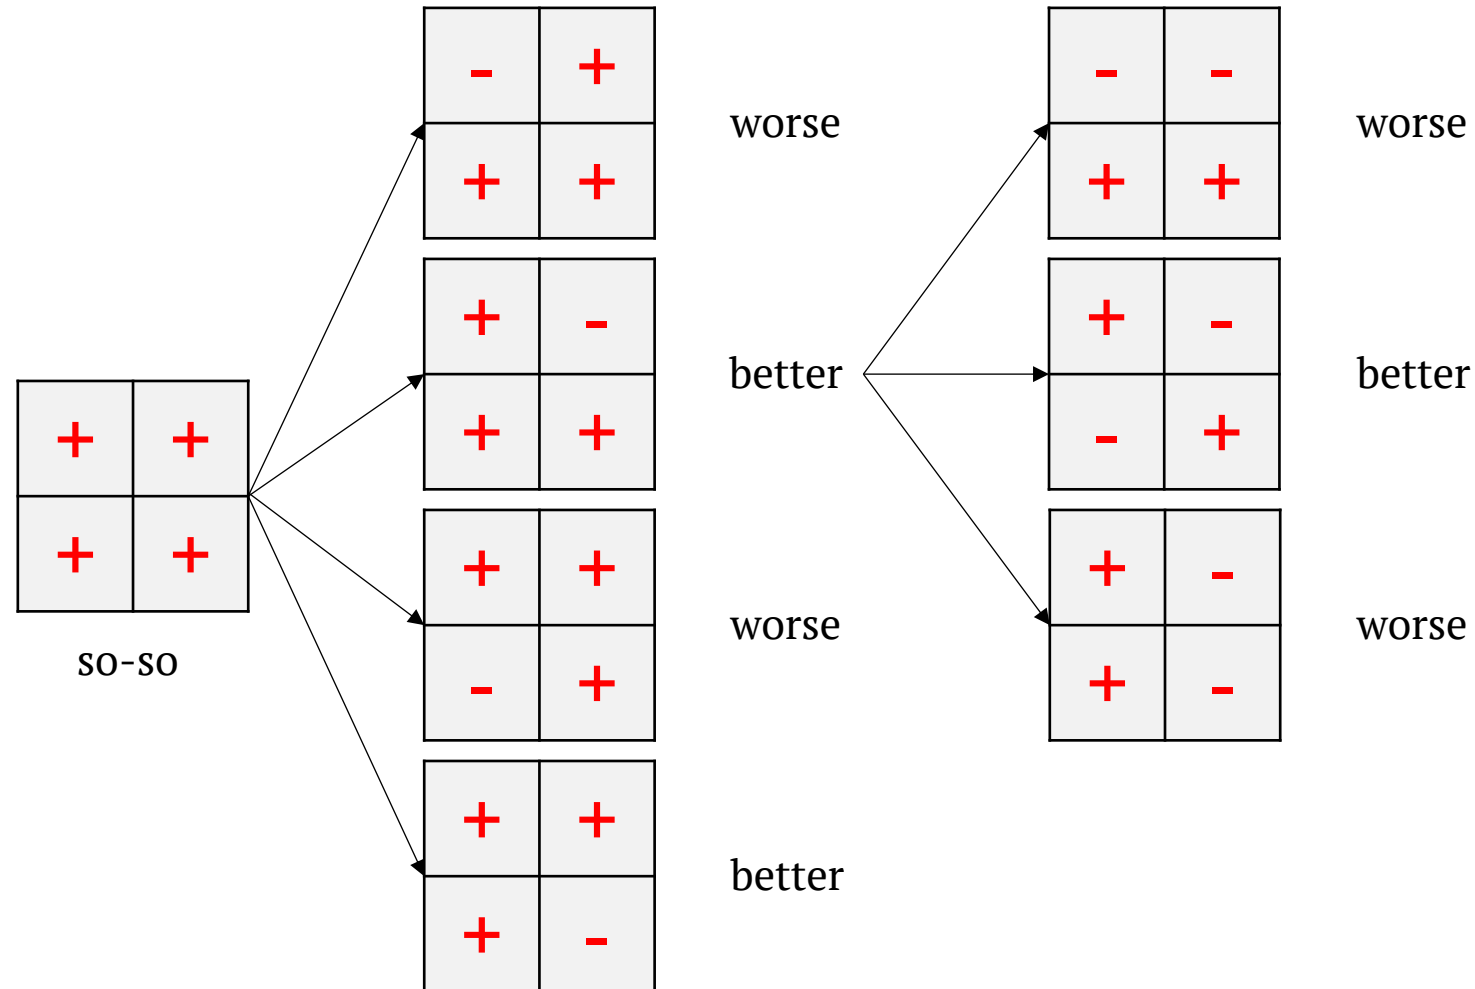
# How can a computer find the filters?

|   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| <table><tr><td>+</td><td>+</td></tr><tr><td>+</td><td>+</td></tr></table> | + | + | + | + | <table><tr><td>-</td><td>+</td></tr><tr><td>+</td><td>+</td></tr></table> | - | + | + | + | <table><tr><td>+</td><td>-</td></tr><tr><td>+</td><td>+</td></tr></table> | + | - | + | + | <table><tr><td>+</td><td>+</td></tr><tr><td>-</td><td>+</td></tr></table> | + | + | - | + |
| +   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| +   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| -   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| +   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| +   | - |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| +   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| +   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| -   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| <table><tr><td>+</td><td>+</td></tr><tr><td>+</td><td>-</td></tr></table> | + | + | + | - | <table><tr><td>-</td><td>-</td></tr><tr><td>+</td><td>+</td></tr></table> | - | - | + | + | <table><tr><td>-</td><td>+</td></tr><tr><td>-</td><td>+</td></tr></table> | - | + | - | + | <table><tr><td>+</td><td>+</td></tr><tr><td>-</td><td>-</td></tr></table> | + | + | - | - |
| +   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| +   | - |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| -   | - |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| +   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| -   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| -   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| +   | + |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |
| -   | - |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |   |

and more ...

Find the best filters from the 16 choices

# How can a computer find the filters?





# How can a computer find the filters?

|     |     |
|-----|-----|
| 0.5 | 1.2 |
| 0.7 | 1.0 |

|     |     |
|-----|-----|
| 0.6 | 0.9 |
| 0.5 | 1.1 |

|      |      |
|------|------|
| 0.9  | -0.7 |
| -0.6 | 1.1  |

and so on...

# Gradient Descent

|     |     |
|-----|-----|
| 0.5 | 1.2 |
| 0.7 | 1.0 |

|     |     |
|-----|-----|
| 0.6 | 0.9 |
| 0.5 | 1.1 |

|      |      |
|------|------|
| 0.9  | -0.7 |
| -0.6 | 1.1  |









|    |    |
|----|----|
| 1  | -1 |
| -1 | 1  |

**Derivatives!**

**Lots of errors**

**Few errors**


# Image Recognition: More Complex Example




# How can a computer find the filters?

|                 |                 |                 |
|-----------------|-----------------|-----------------|
| <sup>+</sup> 1  | <sup>-</sup> -1 | <sup>+</sup> 1  |
| <sup>-</sup> -1 | <sup>+</sup> 1  | <sup>-</sup> -1 |
| <sup>+</sup> 1  | <sup>-</sup> -1 | <sup>+</sup> 1  |

# Previous Knowledge



# Convolutional Neural Network

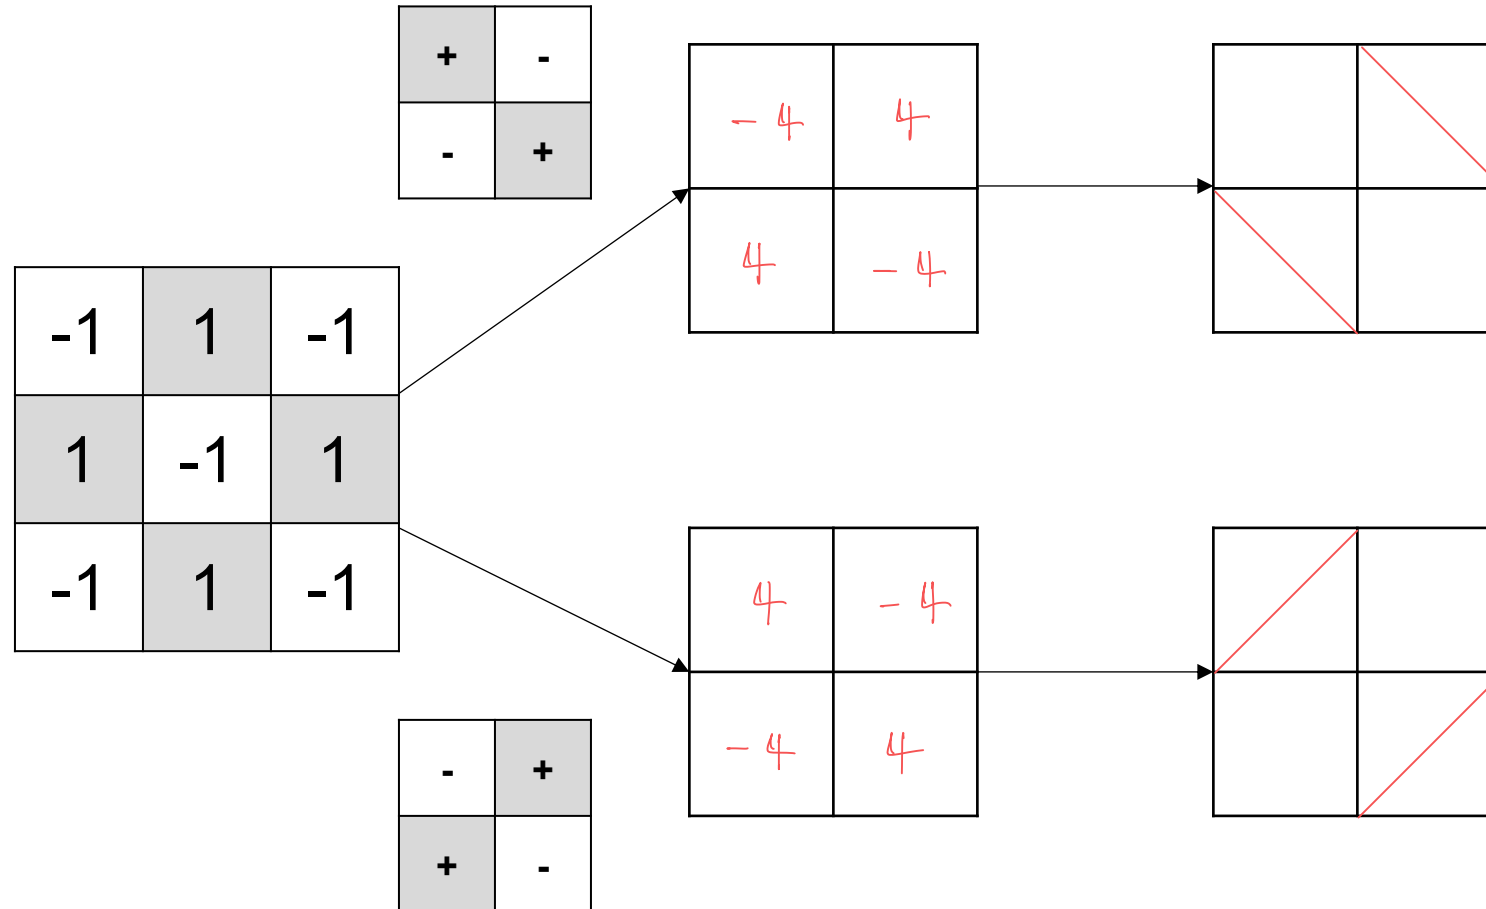


Convolutional Layer  
Pooling Layer

Fully Connected Layer




# Convolutional Neural Network






# Convolutional Neural Network




# Convolutional Neural Network




# Convolutional Neural Network

## Filters




# Convolutional Neural Network




Convolutional Layer  
Pooling Layer

Fully Connected Layer



# Gradient Descent


Lots of errors



Few errors


# Typical CNN architecture

convolutional: filter + input  
pooling: input



Source: mathworks.com

# Input Image





## 2 Components of CNNs

- **Feature extraction** – the hidden layers
  - Convolution layers – the kernel
  - Pooling layers
- **Classification** – the fully connected layers

# Feature Extraction: Convolution Layer

|                 |                 |                 |   |   |
|-----------------|-----------------|-----------------|---|---|
| 1 <sub>x1</sub> | 1 <sub>x0</sub> | 1 <sub>x1</sub> | 0 | 0 |
| 0 <sub>x0</sub> | 1 <sub>x1</sub> | 1 <sub>x0</sub> | 1 | 0 |
| 0 <sub>x1</sub> | 0 <sub>x0</sub> | 1 <sub>x1</sub> | 1 | 1 |
| 0               | 0               | 1               | 1 | 0 |
| 0               | 1               | 1               | 0 | 0 |

Image

## Feature Map


|   |   |   |
|---|---|---|
| 4 | 3 | 4 |
| 2 | 4 | 3 |
| 2 | 3 | 4 |

Convolved  
Feature

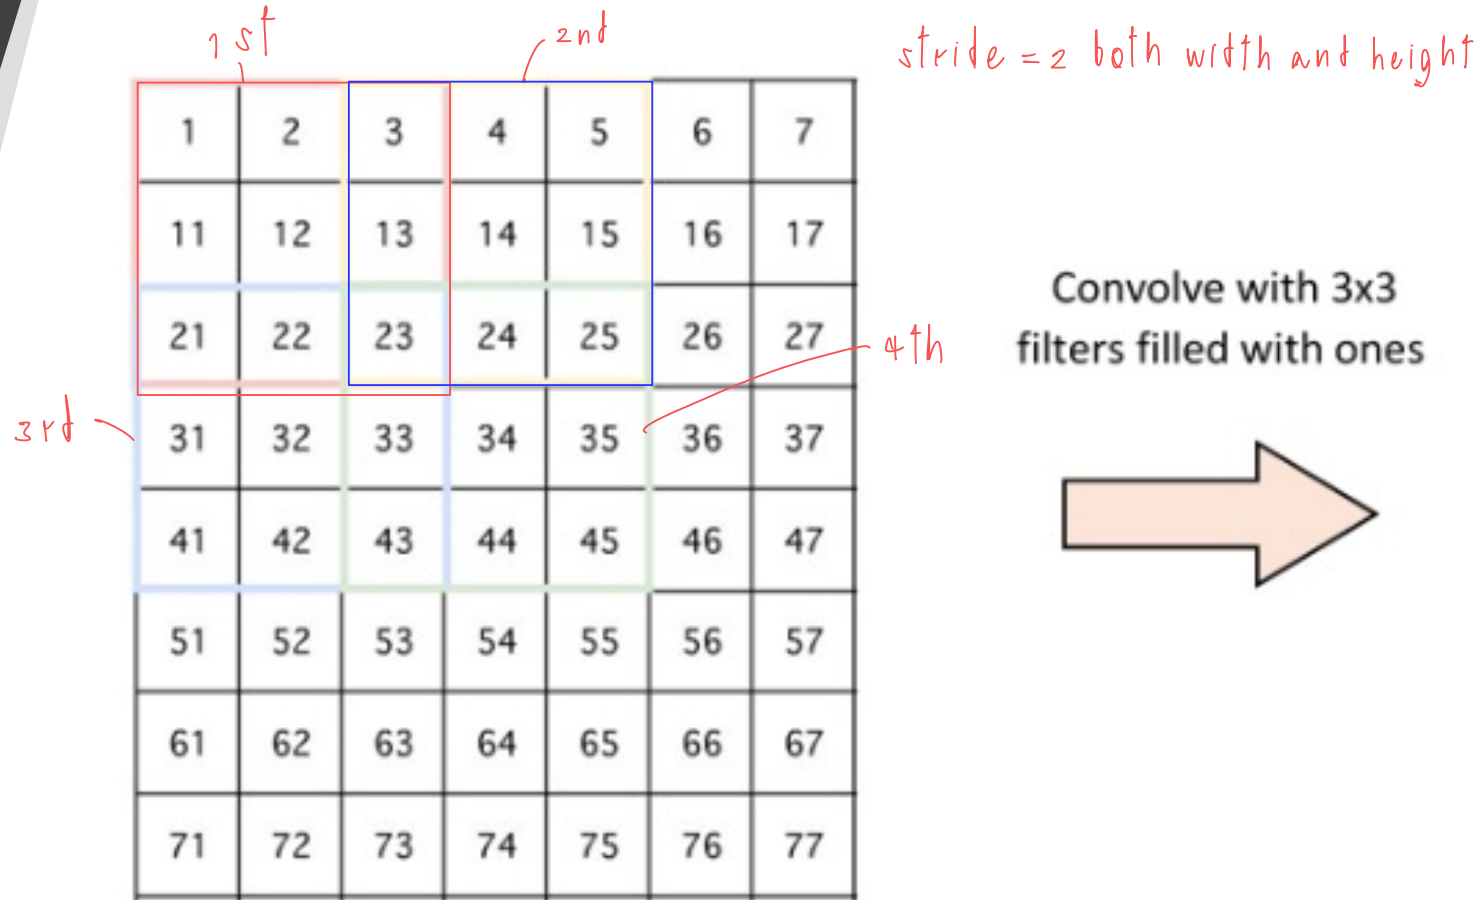
# Convolution Layer: Strides

- It is also possible to connect a large input layer to a much smaller layer by spacing out the receptive fields.
- The distance between two consecutive receptive fields is called the *stride*.

การย้าย convolutional  
matrix ; default = 1



# Convolution Layer: Strides




|     |     |  |
|-----|-----|--|
| 108 | 126 |  |
| 288 | 306 |  |
|     |     |  |

Stride of 2 pixels  
(Source: Raghav Prabhu)

# Convolution Layer: Border Effects and Padding


- In order for a layer to have the same height and width as the previous layer, it is common to add zeros around the inputs, as shown in the diagram. This is called *zero padding*.



# Feature Extraction: Pooling Layer

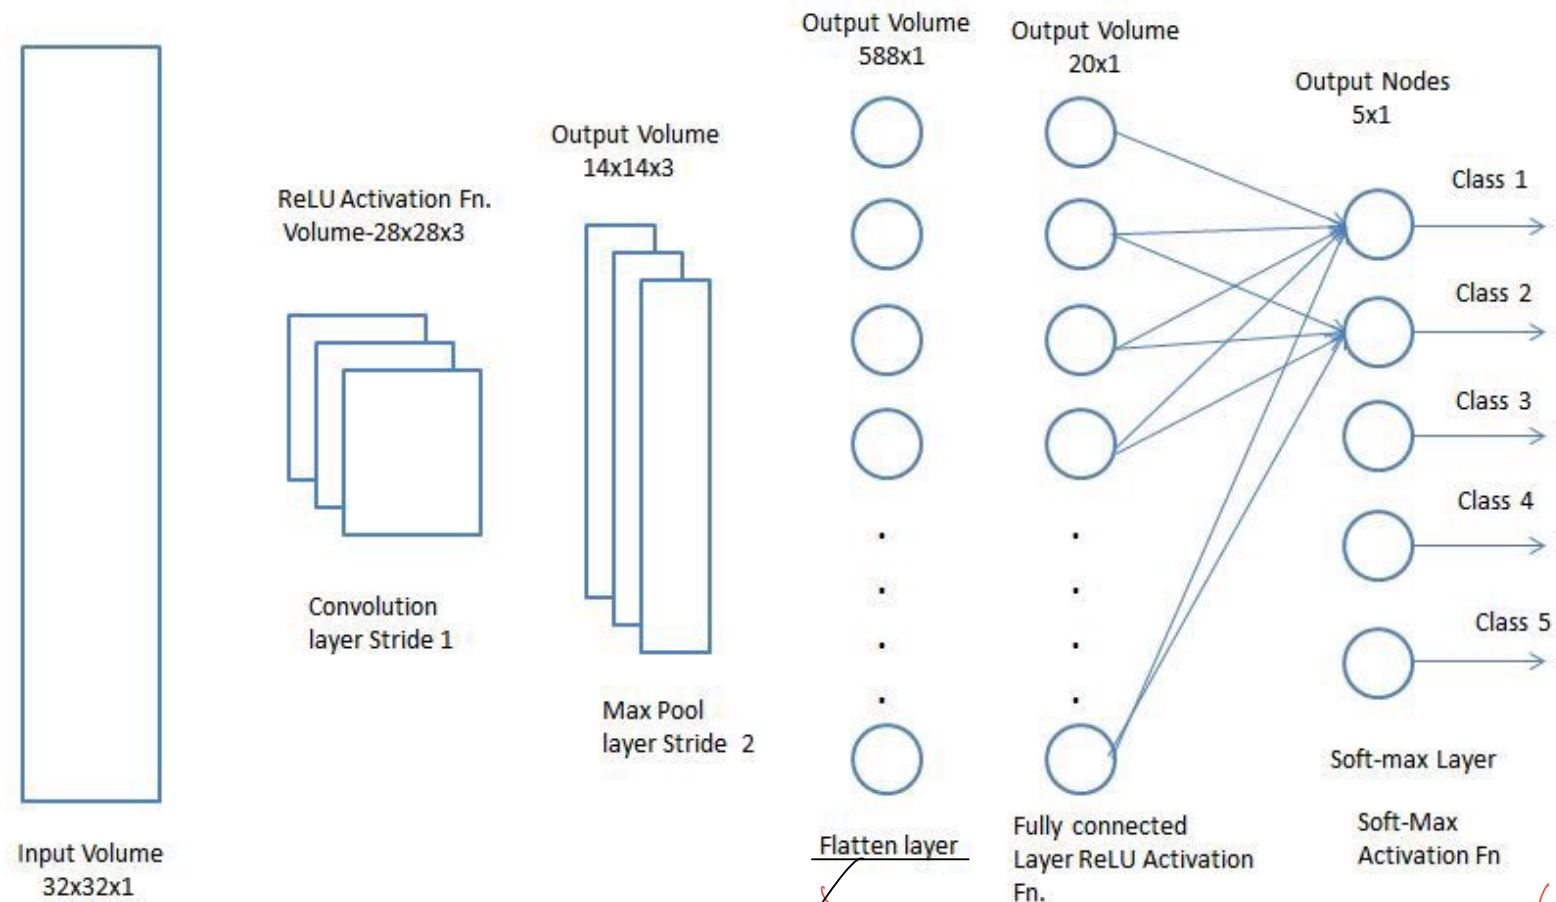
- It is common to add a pooling layer in between CNN layers
  - to continuously reduce the dimensionality
    - to reduce the number of parameters and computation in the network.
  - to shortens the training time
  - to control overfitting.
- Types of pooling
  - Max pooling
  - Average pooling
  - Sum pooling

# Feature Extraction: Pooling Layer



Source: Sumit Saha

# Classification: Fully Connected Layer



จับคู่ข้อมูล  
จากภาพ convolution  
เสร็จ

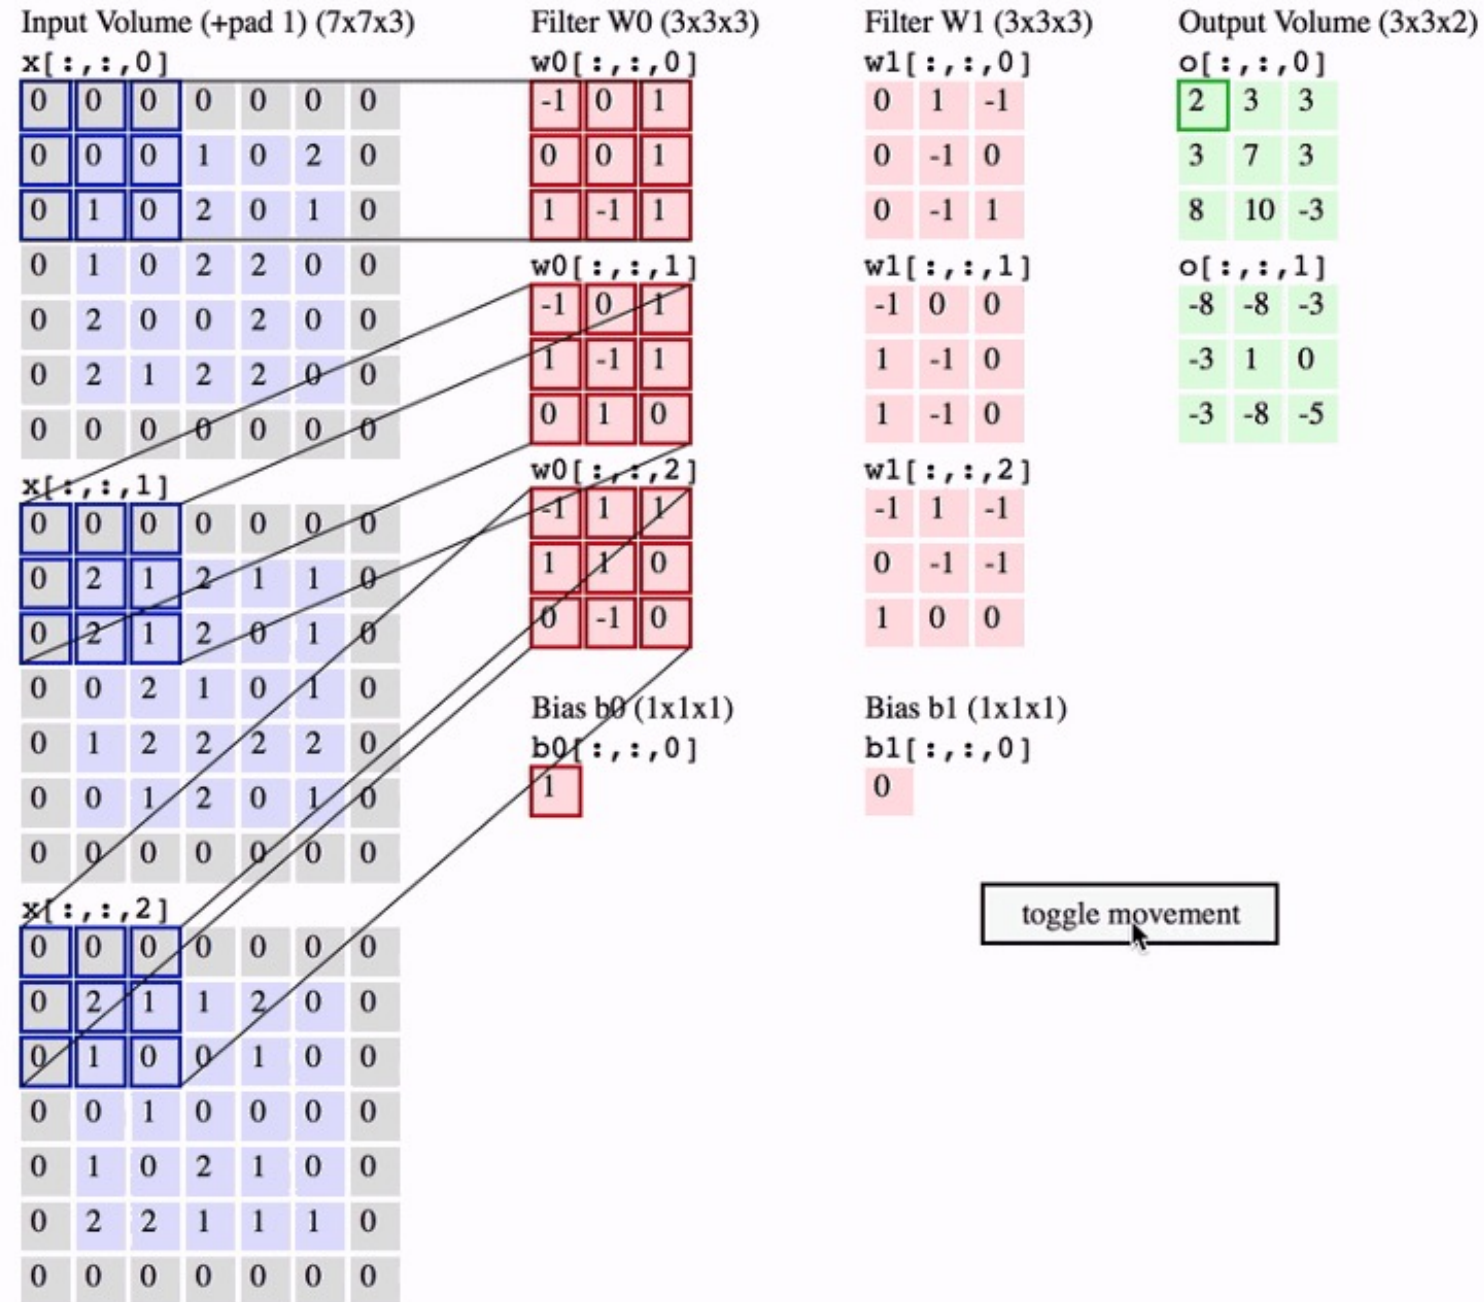
classification




# CNN's 4 Key Hyperparameters

matrix convolution mat.

- The **kernel size**
- The **filter count** (how many filters we want to use)
- **Stride** (how big the steps of the filter are)
- **Padding**



# How convolution works




Note that the output width and height may differ from the input width and height.

They may differ for two reasons:

- **Border effects**, which can be countered by padding the input feature map
- The use of *strides*

# Convolution layer with multiple feature maps



# Other CNN Architectures

- Classical architecture:
  - LeNet-5 (1998)
- Three winners of the **ILSVRC** challenge:
  - AlexNet (2012)
  - GoogLeNet (2014)
  - ResNet (2015)

# LeNet-5 Architecture


|          | Layer | Type            | Maps | Size           | Kernel size  | Stride | Activation |
|----------|-------|-----------------|------|----------------|--------------|--------|------------|
| Classify | Out   | Fully Connected | —    | 10             | —            | —      | RBF        |
|          | F6    | Fully Connected | —    | 84             | —            | —      | tanh       |
| Hidden   | C5    | Convolution     | 120  | $1 \times 1$   | $5 \times 5$ | 1      | tanh       |
|          | S4    | Avg Pooling     | 16   | $5 \times 5$   | $2 \times 2$ | 2      | tanh       |
|          | C3    | Convolution     | 16   | $10 \times 10$ | $5 \times 5$ | 1      | tanh       |
|          | S2    | Avg Pooling     | 6    | $14 \times 14$ | $2 \times 2$ | 2      | tanh       |
|          | C1    | Convolution     | 6    | $28 \times 28$ | $5 \times 5$ | 1      | tanh       |
|          | In    | Input           | 1    | $32 \times 32$ | —            | —      | —          |



# AlexNet

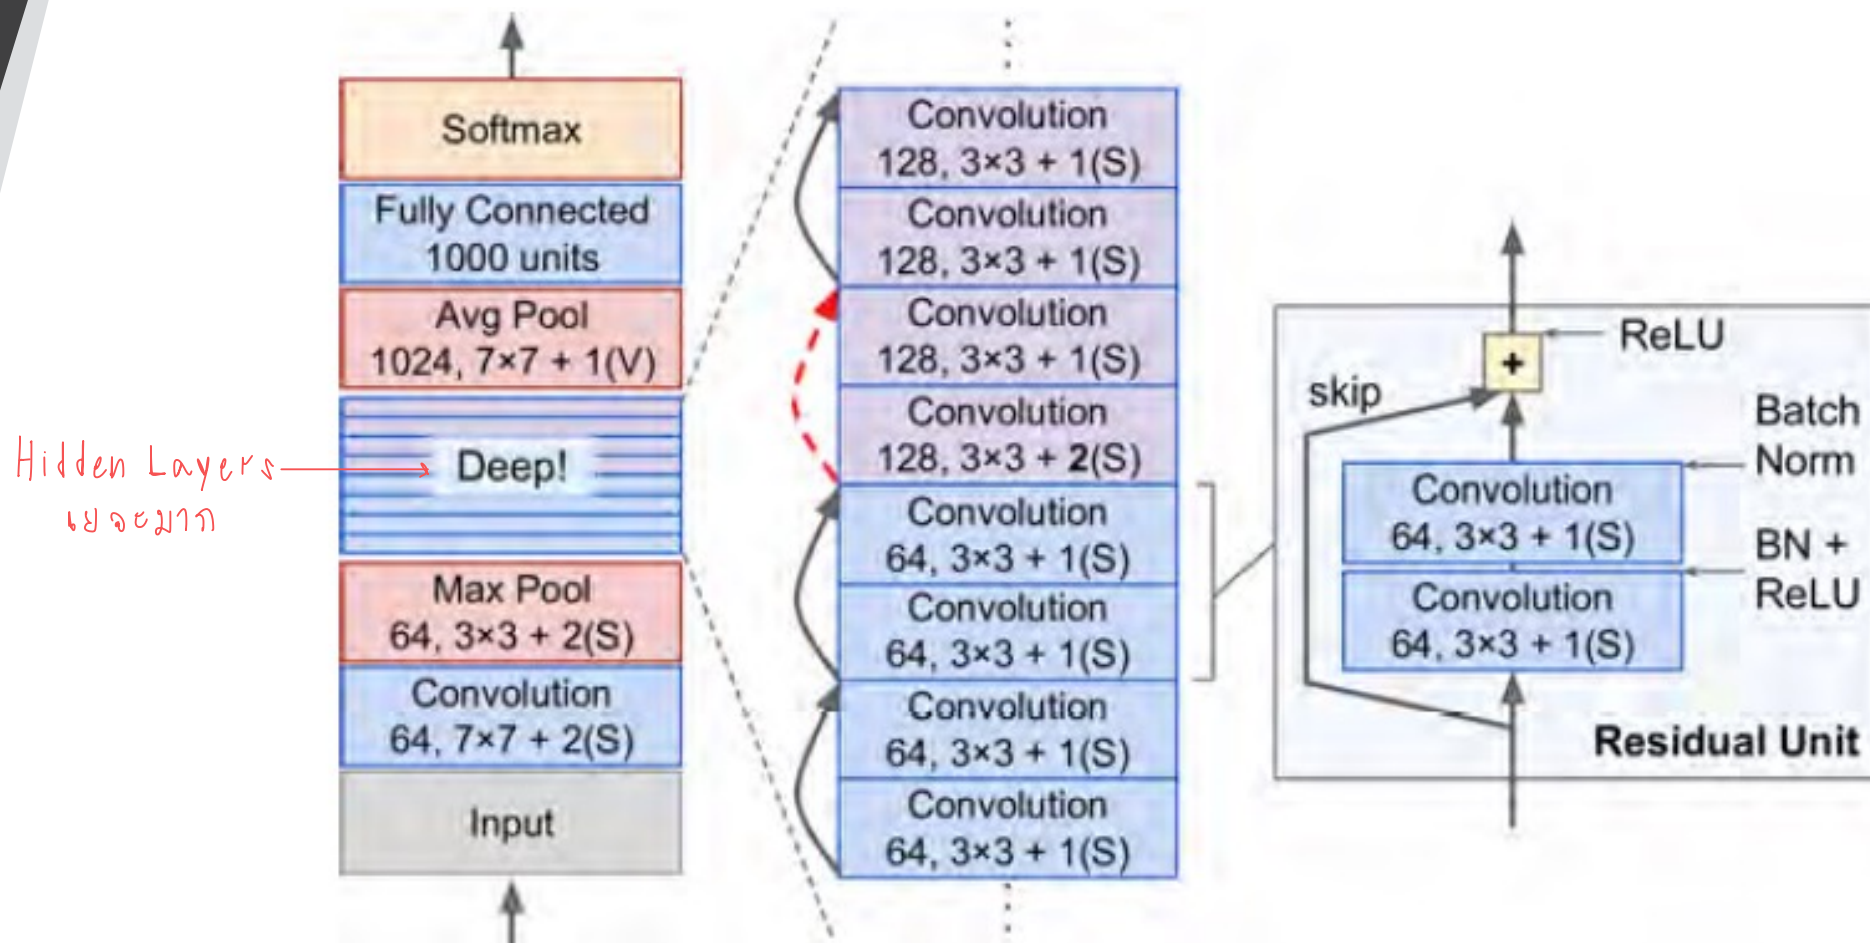
| Layer | Type            | Maps    | Size             | Kernel size    | Stride | Padding | Activation |
|-------|-----------------|---------|------------------|----------------|--------|---------|------------|
| Out   | Fully Connected | —       | 1,000            | —              | —      | —       | Softmax    |
| F9    | Fully Connected | —       | 4,096            | —              | —      | —       | ReLU       |
| F8    | Fully Connected | —       | 4,096            | —              | —      | —       | ReLU       |
| C7    | Convolution     | 256     | $13 \times 13$   | $3 \times 3$   | 1      | SAME    | ReLU       |
| C6    | Convolution     | 384     | $13 \times 13$   | $3 \times 3$   | 1      | SAME    | ReLU       |
| C5    | Convolution     | 384     | $13 \times 13$   | $3 \times 3$   | 1      | SAME    | ReLU       |
| S4    | Max Pooling     | 256     | $13 \times 13$   | $3 \times 3$   | 2      | VALID   | —          |
| C3    | Convolution     | 256     | $27 \times 27$   | $5 \times 5$   | 1      | SAME    | ReLU       |
| S2    | Max Pooling     | 96      | $27 \times 27$   | $3 \times 3$   | 2      | VALID   | —          |
| C1    | Convolution     | 96      | $55 \times 55$   | $11 \times 11$ | 4      | SAME    | ReLU       |
| In    | Input           | 3 (RGB) | $224 \times 224$ | —              | —      | —       | —          |

# GoogLeNet





# ResNet Architecture



# Transfer Learning



"I think transfer learning is the key to general intelligence. And I think the key to doing transfer learning will be the acquisition of conceptual knowledge that is abstracted away from perceptual details of where you learned it from."

- Demis Hassabis  
CEO, DeepMind

# What is transfer learning?

- A machine learning technique where a **model trained** on one task is re-purposed on a **second related task**
  - For example, if you trained a simple classifier to predict whether an image contains a backpack, you could use the knowledge that the model gained during its training to recognize other objects like sunglasses.
- Mostly used in Computer Vision and Natural Language Processing Tasks
  - because of the huge amount of computational power that is needed for them

# How to use transfer learning?

- Two common approaches:
  - Develop model → ของตนเอง
  - Pre-trained model → เอาของคนอื่นมา

# Develop Model Approach

## 1. Select Source Task.

- select a related predictive modeling problem with an abundance of data

## 2. Develop Source Model.

- develop a skillful model for this first task
- The model must be better than a naive model

## 3. Reuse Model.

- The model fit on the source task can then be used as the starting point for a model on the second task of interest.
- This may involve using **all or parts of the model**, depending on the modeling technique used.

## 4. Tune Model.

# Pre-Trained Model Approach

## 1. Select **Source Model**.

- choose an available pre-trained source model

## 2. Reuse Model.

- use the pre-trained model can as the starting point for the second task of interest

## 3. Tune Model.

\* common in the field of deep learning \*

# Examples of Transfer Learning with Image Data

- It is common to use a deep learning model pre-trained for a large and challenging image classification task such as the **ImageNet** 1000-class photograph classification competition
- The research organizations often release their final model under a permissive license for reuse
  - **Oxford VGG Model:**  
[http://www.robots.ox.ac.uk/~vgg/research/very\\_deep/](http://www.robots.ox.ac.uk/~vgg/research/very_deep/)
  - **Google Inception Model**  
<https://github.com/tensorflow/models/tree/master/inception>
  - **Microsoft ResNet Model**  
<https://github.com/KaimingHe/deep-residual-networks>
- These models can take days or weeks to train on modern hardware.
- These models can be downloaded and incorporated directly into new models that expect image data as input.

# Training CNN in Keras



```
cnn = models.Sequential()
cnn.add(layers.Conv2D(40, kernel_size=5, padding="same",
                      input_shape=(28, 28, 1), activation = 'relu', name = 'conv1_1'))
cnn.add(layers.MaxPool2D((2, 2)))
cnn.add(layers.Dropout(0.25))
cnn.add(layers.Conv2D(32, kernel_size=(3, 3),
                      activation='relu', kernel_initializer='he_normal', name = 'conv1_2'))
cnn.add(layers.MaxPool2D((2, 2)))
cnn.add(layers.Dropout(0.25))
cnn.add(layers.Flatten())
cnn.add(layers.Dense(64, activation='relu'))
cnn.add(layers.BatchNormalization())
cnn.add(layers.Dropout(0.25))
cnn.add(layers.Dense(10, activation='softmax'))
```



**Input Mnist data**

**Shape (28 , 28 , 1)**

```
cnn = models.Sequential()
cnn.add(layers.Conv2D(40, kernel_size=5, padding="same",
                      input_shape=(28, 28, 1), activation = 'relu',name = 'conv1_1'))
cnn.add(layers.MaxPool2D((2, 2)))
cnn.add(layers.Dropout(0.25))
cnn.add(layers.Conv2D(32, kernel_size=(3, 3),
                      activation='relu',kernel_initializer='he_normal',name = 'conv1_2'))
cnn.add(layers.MaxPool2D((2, 2)))
cnn.add(layers.Dropout(0.25))
cnn.add(layers.Flatten())
cnn.add(layers.Dense(64, activation='relu'))
cnn.add(layers.BatchNormalization())
cnn.add(layers.Dropout(0.25))
cnn.add(layers.Dense(10, activation='softmax'))
```

Covolution layer  
40, 5\*5

MaxPooling 2\*2  
Dropout(0.25)

Covolution layer  
32, 3\*3

MaxPooling 2\*2  
Dropout(0.25)

Flatten layer

Hidden layer  
64 Node

BatchNormalization()  
Dropout(0.25)

Output layer  
10 Node

| Layer (type)                                | Output Shape       | Param # |
|---|--------------------|---------|
| conv1_1 (Conv2D)                            | (None, 28, 28, 40) | 1040    |
| max_pooling2d_1 (MaxPooling2)               | (None, 14, 14, 40) | 0       |
| dropout_1 (Dropout)                         | (None, 14, 14, 40) | 0       |
| conv1_2 (Conv2D)                            | (None, 12, 12, 32) | 11552   |
| max_pooling2d_2 (MaxPooling2)               | (None, 6, 6, 32)   | 0       |
| dropout_2 (Dropout)                         | (None, 6, 6, 32)   | 0       |
| flatten_1 (Flatten)                         | (None, 1152)       | 0       |
| dense_1 (Dense)                             | (None, 64)         | 73792   |
| batch_normalization_1 (Batch Normalization) | (None, 64)         | 256     |
| dropout_3 (Dropout)                         | (None, 64)         | 0       |
| dense_2 (Dense)                             | (None, 10)         | 650     |
| Total params: 87,290                        |                    |         |
| Trainable params: 87,162                    |                    |         |
| Non-trainable params: 128                   |                    |         |

# Data augmentation

```
# Define a generator for train set and test set

train_datagen = image.ImageDataGenerator(rescale=1./255,
                                         rotation_range=40,
                                         width_shift_range=0.2,
                                         height_shift_range=0.2,
                                         shear_range=0.2,
                                         zoom_range=0.2,
                                         horizontal_flip=False)

test_datagen = image.ImageDataGenerator(rescale=1./255)
```


```
# Create an Iterator object.
train_generator = train_datagen.flow(X_train,y_train,
                                     batch_size = BATCH_SIZE,
                                     seed=0)

validate_generator = test_datagen.flow(X_val,y_val,
                                       batch_size = BATCH_SIZE,
                                       shuffle=False)
```

Using the ImageDataGenerator module to generate more data.

It help generate more variation of the data which help prevent overfit and generalize better.

<https://keras.io/preprocessing/image>





# Transfer Learning

```
from keras.applications import vgg16

vgg = vgg16.VGG16(include_top=False,
                  weights='imagenet',
                  input_shape=(150,150,3))

prev_cnn = models.load_model('your_previos_model.h5')
prev_cnn.summary()
```

# Use .pop() to remove the last layer  
# In this case, we want to remove last two layer

```
prev_cnn.pop()
prev_cnn.pop()
```

If we don't want to train these layer, we have to freeze these layer.

```
prev_cnn.trainable = False
```

Or Freeze a specific layers

# Freeze first 3 layer → ห้ามยุ่ง / ห้ามปรับ weight

```
for i in range(3):
    prev_cnn.layers[i].trainable = False
```

What is transfer learning <https://towardsdatascience.com/transfer-learning-946518f95666>  
which transfer learning method to use <https://medium.com/@14prakash/transfer-learning-using-keras-d804b2e04ef8>  
The following is a tutorial code to load and freeze some layer